

Section 5.4

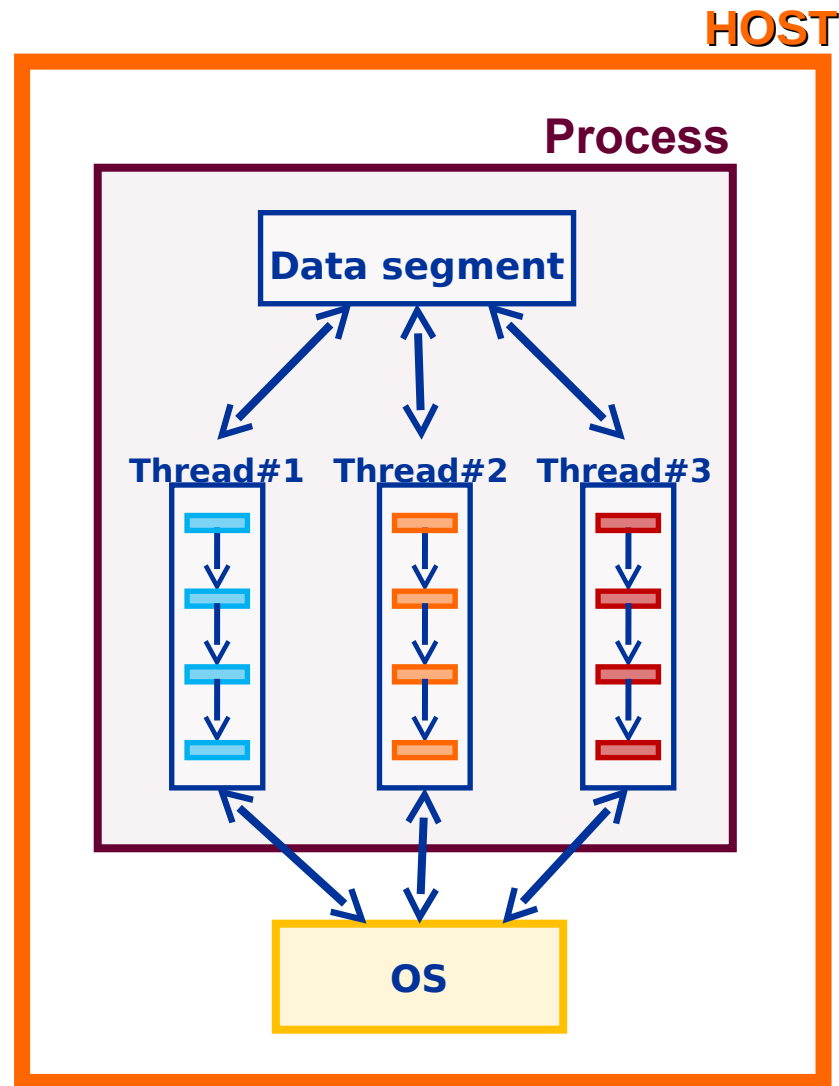
Threads

1. Overview
2. Sharing data

5.4.1 Overview

- What is a *thread*?
 - it's a single logical flow of control inside a process
- Thread management
 - threads can be created by the main thread or by peer threads
 - thread creation specifies a function for the thread to execute
 - threads run in parallel
 - they are scheduled automatically by the OS kernel
 - the context switch between threads is faster than between processes
- **coding example <p1>**

Overview (cont.)



Overview (cont.)

- Characteristics of threads
 - each thread has its own unique resources (*thread context*):
 - a thread id
 - a function call stack
 - a program counter
 - points to instruction currently executing
 - all threads within a process share the same:
 - address space
 - data segment
 - including the heap
 - code segment
- coding example <p2>

5.4.2 Sharing Data

- What areas of memory are shared between threads?
 - global memory
 - heap
- Problem
 - two threads modifying the same data at the same time
- **coding example <p3>**

Sharing Data (cont.)

- Solution:
 - protect the shared data
 - ensure that changes to the shared data are **atomic**
 - atomic operations are *non-divisible* operations
 - their time-slice on the CPU cannot be *preempted*
 - it cannot be interrupted by another thread or process
 - mechanisms to protect shared data:
 - semaphore
 - mutex

Semaphores

- What is a *semaphore*?
 - it's a variable used in only *two* atomic operations
 - *lock* operation
 - *unlock* operation
 - it's initialized with the number of threads that can lock it simultaneously
 - if it's locked, other threads trying to access it are blocked
 - once unlocked, *no guarantee* which thread gets unblocked

Semaphores (cont.)

- Operations on semaphores
 - wait operation
 - this is the *locking* operation
 - if the semaphore is non-zero, operation decrements the semaphore
 - if the semaphore is zero, the thread is blocked until it's non-zero
 - post operation
 - this is the *unlocking* operation
 - it increments the semaphore

Mutexes

- What is a *mutex*?
 - it's a binary semaphore
 - its value is zero or one
 - only *one thread* can access it at a time
 - mutex == **mut**ual **ex**clusion
- coding examples <p4> and <p5>